

main

March 6, 2024

1 Imports

```
[ ]: # imports
from typing import Tuple, List
from collections import deque

import torch
import torch.optim as optim
from torch.optim import lr_scheduler
from torch.optim import Optimizer

import numpy as np
from torch import Tensor

import torch.nn as nn
import torch.nn.functional as F
from torch.nn.modules.loss import _Loss

torch.manual_seed(20190325);

[ ]: %load_ext autoreload
%autoreload 2

[ ]: import torchvision
from torchvision.datasets import MNIST
import torchvision.transforms as transforms
from torch.utils.data import DataLoader

img_transforms = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.1305,), (0.3081,))
])

[ ]: # https://pytorch.org/docs/stable/data.html
train_dataset = MNIST(root='../mnist_data/',
                       train=True,
                       download=True,
```

```

        transform=img_transforms)

test_dataset = MNIST(root='../mnist_data/',
                     train=False,
                     download=True,
                     transform=img_transforms)

```

```

[ ]: train_loader = torch.utils.data.DataLoader(dataset=train_dataset,
                                                batch_size=60,
                                                shuffle=True)

test_loader = torch.utils.data.DataLoader(dataset=test_dataset,
                                           batch_size=60,
                                           shuffle=False)

```

```

[ ]: from model import Autoencoder
     from train import PyTorchTrainer

```

2 Unsupervised Learning with Autoencoder

2.1 Preprocess data

```

[ ]: mnist_train = ((train_dataset.data.type(torch.float32).unsqueeze(3).permute(0, ↵
↵3, 1, 2) / 255.0) - 0.1305) / 0.3081
mnist_test = ((test_dataset.data.type(torch.float32).unsqueeze(3).permute(0, 3, ↵
↵1, 2) / 255.0) - 0.1305) / 0.3081

```

```

[ ]: X_train = mnist_train
     X_test = mnist_test

```

```

[ ]: X_train_auto = (X_train - X_train.min()) / (X_train.max() - X_train.min()) * 2 ↵
↵- 1
     X_test_auto = (X_test - X_train.min()) / (X_train.max() - X_train.min()) * 2 - 1

```

```

[ ]: model = Autoencoder()

```

```

[ ]: model = Autoencoder(hidden_dim=28)
     criterion = nn.MSELoss()
     optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.9)

     trainer = PyTorchTrainer(model, optimizer, criterion)

     trainer.fit(X_train_auto, X_train_auto,
                 X_test_auto, X_test_auto,
                 epochs=1,
                 batch_size=60)

```

The loss after 1 epochs was 0.06897220760583878

```
[ ]: reconstructed_images, image_representations = model(X_test_auto)
```

```
[ ]: import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
[ ]: def display_image(ax,
    t: Tensor):
    n = t.detach().numpy()
    ax.imshow(n.reshape(28, 28))
```

```
[ ]: np.random.seed(20190504)
a = np.random.randint(0, 10000)
```

```
[ ]: X_test[a].shape
```

```
[ ]: torch.Size([1, 28, 28])
```

```
[ ]: f, axarr = plt.subplots(1,2)
display_image(axarr[0], X_test[a])
display_image(axarr[1], reconstructed_images[a])

axarr[0].set_title("Original image")
axarr[1].set_title("Image reconstructed\nfrom autoencoder")

axarr[0].axis('off')
axarr[1].axis('off');

# f.savefig("../01_deep-learning-from-scratch/images/07_pytorch/
↪03_autoencoder_example_image.png")
```

Original image





Image reconstructed
from autoencoder



3 t -SNE on the results

```
[ ]: from sklearn.manifold import TSNE
tsne_result = TSNE(n_components=2, random_state=20190405).
    ↪ fit_transform(image_representations.detach().numpy())
```

3.1 t -SNE viz

```
[ ]: import pandas as pd
tsne_df = pd.DataFrame({'tsne_dim_1': tsne_result[:,0],
                        'tsne_dim_2': tsne_result[:,1],
                        'category': test_dataset.targets})
groups = tsne_df.groupby('category')

# Plot
fig, ax = plt.subplots(figsize=(25,25))
ax.set_title('10000 observations from MNIST test set, colored by their actual_
    ↪ digit.
Locations are the result of reducing the 28 values from hidden layer of the_
    ↪ convolutional
autoencoder - trained without labels - down to two dimensions using t-SNE.')
ax.margins(0.05) # Optional, just adds 5% padding to the autoscaling
for name, group in groups:
    ax.scatter(group['tsne_dim_1'], group['tsne_dim_2'], marker='o', label=name)
ax.legend();
# fig.savefig("../01_deep-learning-from-scratch/images/07_pytorch/00_tsne.
    ↪ png")
```

